

# CLOP-DiT: Text-Guided Single-Cell Gene Expression Generation via Contrastive Language-Omics Pretraining and Diffusion Transformers

Zeyu Fu <sup>1,\*</sup> 

<sup>1</sup> State Key Laboratory of Trauma and Chemical Poisoning, Institute of Combined Injury, Chongqing Engineering Research Center for Nanomedicine, College of Preventive Medicine, Army Medical University, Chongqing 400038, China; zeyu.fu@tmmu.edu.cn

## Abstract

Generating realistic synthetic single-cell gene expression profiles conditioned on natural-language descriptions of cell identity would enable data augmentation and rare cell-type simulation, and could in future support hypothesis-driven perturbation studies. We present CLOP-DiT, a two-stage pipeline combining Contrastive Language–Omics Pretraining (CLOP) with a conditional Diffusion Transformer (DiT) trained via flow matching. CLOP aligns BiomedBERT text embeddings and scGPT cell embeddings into a shared 512-dimensional space (3.02M parameters), amplifying inter-type separation from 0.6% to 78% in cosine distance. DiT1D (22.10M parameters, 8 AdaLN-Zero blocks) then performs conditional flow matching in scGPT latent space, and a frozen scGPT decoder maps generated latents to gene expression. Across 100 cell types from 80 datasets (220,304 cells, 1,088 text groups), CLOP-DiT achieves  $k$ -nearest-neighbor accuracy of 36.9% ( $37\times$  random chance, but 52 percentage points below the 89% real-data baseline) and steering accuracy of 81.0% (random = 50%) in the high-fidelity regime (CFG = 2.0, Euler solver); the recommended high-diversity regime (CFG = 1.0, Midpoint solver) reaches a diversity ratio of 0.93 at 80.7% steering. Unconditional generation collapses to random baselines, confirming that text conditioning is the sole driver of type specificity. In-distribution downstream biological validation—clustering alignment, classifier transfer (expression-space logistic regression accuracy 30.8%, macro F1 0.247), and differential expression concordance (mean logFC Pearson  $r \approx 0.39$  across three contrasts, range 0.17–0.56)—demonstrates that generated cells partially integrate into standard single-cell workflows, though a discriminator AUC of 0.656 and the 52-point KNN gap from real data indicate substantial room for improvement.

**Keywords:** single-cell RNA-seq; generative model; diffusion transformer; contrastive learning; text-guided generation; flow matching; gene expression; foundation model

## 1. Introduction

Single-cell RNA sequencing (scRNA-seq) has enabled the characterization of cellular heterogeneity at unprecedented resolution [1,2]. However, generating realistic synthetic single-cell expression profiles conditioned on natural-language descriptions of cell identity remains an open challenge. Such a capability would enable data augmentation for rare cell types, and could potentially support hypothesis testing via *in silico* perturbation and atlas completion without additional sequencing, though these applications remain to be validated.

Received:

Revised:

Accepted:

Published:

**Copyright:** © 2026 by the author.

Submitted to *Biology* for possible open access publication under the terms and conditions of the [Creative Commons Attribution \(CC BY\)](https://creativecommons.org/licenses/by/4.0/) license.

The clinical motivation for text-conditioned generation is threefold. First, rare cell populations—circulating tumor cells, therapy-resistant subclones, and transitional progenitor states—are chronically undersampled in standard scRNA-seq experiments, yet their accurate representation is critical for biomarker discovery and treatment-response prediction. Synthetic augmentation of these populations could improve classifier balance and statistical power without additional sequencing cost. Second, large-scale atlas efforts such as the Human Cell Atlas [3] require imputation of missing cell states across tissue contexts; a generative model conditioned on descriptive text could fill gaps where donor material is unavailable. Third, *in silico* hypothesis generation—predicting the transcriptomic consequences of a cell-state transition described in natural language—could accelerate experimental design in perturbation biology, although this application remains speculative and unvalidated.

A fundamental technical gap separates current generative methods from this vision. Variational autoencoders [4] such as scVI [5,6] and scGen [7], geometric deep generative models [8], and foundation models such as Geneformer [9] and scBERT [10–13] all condition on categorical cell-type labels or perturbation metadata. To our knowledge, no existing method accepts structured natural-language biological descriptions as conditioning input for single-cell generation, though related text-conditioned approaches for cell-level analysis have been explored [49,50]. In parallel, contrastive vision–language models such as CLIP [14] have demonstrated that aligning two modalities into a shared embedding space enables powerful zero-shot transfer, with SimCLR [15] and MoCo [16] establishing the self-supervised contrastive paradigm and domain-specific variants such as BiomedCLIP [17] and SigLIP [18] extending it to biomedical data. Meanwhile, Diffusion Transformers (DiTs) [19] have emerged as state-of-the-art conditional generators, building on denoising diffusion probabilistic models [20,21] and latent diffusion models [22–24], and flow matching [25,51,54] provides a simulation-free training objective that simplifies the diffusion framework. The Transformer architecture [26] underpins both the DiT generator and the frozen encoders. Bridging these two lines of work—contrastive language–omics alignment and conditional diffusion generation—is the central contribution of this paper.


We present CLOP-DiT, a two-stage pipeline that combines **C**ontrastive **L**anguage–**O**mic **P**retraining (CLOP) with a conditional **D**iffusion **T**ransformer (DiT) to generate single-cell gene expression profiles from text descriptions. The pipeline (Figure 1) first aligns text and cell embeddings into a shared contrastive space, then generates latents via flow matching, and finally decodes to gene expression. Building on BiomedBERT [27] and scGPT [28], CLOP-DiT is, to our knowledge, the first method to combine contrastive language–omics alignment with a diffusion transformer for text-conditioned single-cell generation. Across 100 evaluated cell types, CLOP-DiT achieves a KNN classification accuracy of 36.9% ( $37\times$  random chance, but 52 points below the 89% real-data ceiling), a steering accuracy of 81.0%, and a diversity ratio of 0.93, indicating that text conditioning drives detectable type-specific generation, albeit with a substantial fidelity gap from real data.

The principal contributions of this work are: (1) a contrastive alignment module (CLOP) that maps text and cell embeddings into a shared space via PrototypeSigLIP loss, amplifying inter-group cosine separation from 0.006 to 0.778; (2) a 1D Diffusion Transformer (DiT1D, 22.1M parameters) that performs conditional flow matching in the 512-dimensional scGPT latent space; (3) a comprehensive four-axis evaluation framework—KNN fidelity, steering accuracy, diversity ratio, and downstream biological validation—applied to 100 cell types with bootstrap confidence intervals; and (4) a cross-dataset biological validation pipeline demonstrating near-perfect gene-level fidelity (Pearson  $r > 0.999$ ) across five independent tissue contexts.

The remainder of this paper is organized as follows: Sections 2.1–2.5 describe the dataset, model architecture, and evaluation framework; Section 3 presents training dynamics, generation quality, and downstream validation results; Section 4 discusses interpretation, limitations, and applications; and Section 5 summarizes the contributions.

## 2. Materials and Methods

Figure 1 provides an overview of the pipeline. We organize the Materials and Methods in three parts: (1) Dataset curation and preprocessing (2.1); (2) Model architecture and training—covering CLOP alignment (2.2), DiT generation (2.3), and scGPT decoding (2.4); (3) Evaluation framework and metrics (2.5). Together, these components form a reproducible pipeline for text-conditioned single-cell generation.



**Figure 1.** Overview of the CLOP-DiT pipeline. **Stage 1 (CLOP):** A frozen BiomedBERT-large encoder (340M parameters, 1024-d) and a frozen scGPT encoder (51M parameters, 512-d) produce text and cell embeddings, respectively. Dual MLP projectors map both modalities to a shared 512-dimensional  $L_2$ -normalized space, trained with a PrototypeSigLIP contrastive loss. ZCA whitening is applied to BiomedBERT embeddings prior to projection. **Stage 2 (DiT):** A 1D Diffusion Transformer (8 AdaLN-Zero blocks, 22.1M parameters) performs conditional flow matching in the 512-dimensional scGPT latent space, using CLOP-projected text as the conditioning signal. An Euler or Midpoint ODE solver integrates the learned velocity field from Gaussian noise  $z_0 \sim \mathcal{N}(0, I)$  to a generated latent  $z_1$ . Classifier-free guidance (CFG) is applied at inference. **Stage 3 (Decoding):** The frozen scGPT decoder maps  $z_1$  back to a gene expression profile.

### 2.1. Dataset

We curated 80 datasets from the Gene Expression Omnibus (GEO) spanning cancer and developmental biology, comprising 220,304 cells after quality control and highly variable gene selection [29,53] (1,790 genes selected by mean expression and dispersion). Each cell was encoded into a 512-dimensional embedding by a pre-trained scGPT model. Cell-type annotations were organized into 1,088 unique text groups, each described by a structured natural-language caption (not free-form text) incorporating cell type, marker genes, tissue of origin, organism, and disease context. Text descriptions were generated for each cell type using a structured template: “{cell type}, tissue: {tissue}, organism: {organism}, markers: {top 5 DE marker genes}, context: {disease/condition}”, where marker genes were identified per cell type by one-vs.-rest Wilcoxon rank-sum test [58] on the training set. The exact template and marker gene sources are available in `src/data_pipeline/subcluster_annotation` (caption generation) and `scripts/data_prep/02b_enrich_descriptions.py` (evidence enrichment).

The corpus was split into 72 training datasets (199,670 cells) and 8 held-out validation datasets (20,634 cells). The 8 validation datasets were selected by study (not by cell) to ensure no sample overlap between training and validation splits, and were chosen to span

the major tissue and cancer categories in the corpus (lung, gastric, skin, liver, blood) via stratified selection across tissue types; the full list of GEO accession identifiers for all 80 datasets is provided in Supplementary Table S1, and the validation study identifiers are in Supplementary Table S2. Within the 72 training datasets, a cell-type-stratified split reserves ~10% of cells per type for training-time validation, ensuring that all 69 deduplicated cell types are represented in both partitions. This held-out design tests interpolation within the training distribution; the 8 held-out datasets share tissue types and cell populations with the training corpus and do not represent a stringent out-of-distribution generalization test.

The relationship between the 1,088 text groups and the 100 evaluated cell types is as follows. The 1,088 text groups correspond to sub-cluster-level descriptions: each Leiden [30] cluster within each dataset receives a unique text caption. Many sub-clusters share the same canonical cell type (e.g., “CD8+ T cells” in lung and liver datasets produce separate text groups). For evaluation, text groups are consolidated by core cell-type identity, retaining only groups with at least 20 cells (to ensure stable centroid estimation), yielding 100 evaluable cell types. The 100-type, 200-cells-per-group design follows standard single-cell generative evaluation practice [5,7]; the minimum threshold of 20 real cells per type ensures stable centroid estimation. KNN accuracy is therefore computed over these 100 consolidated types (random chance =  $1/100 = 1.0\%$ ). The text–cell alignment heatmap (Figure 5 and 6) uses 69 centroids because it operates on the deduplicated condition space, which merges cross-dataset sub-clusters into canonical types.

The 80 datasets comprise both *Homo sapiens* (~59 datasets) and *Mus musculus* (~21 datasets) [31], predominantly from tumor microenvironment and developmental contexts. Consequently, all generation and evaluation results should be interpreted as applying to human and mouse cancer and developmental single-cell biology; generalization to other organisms, non-cancer tissues, or perturbation conditions has not been tested.

Training and evaluation were implemented in Python 3.10 using PyTorch 2.1 (CUDA 13.0), scGPT v0.2.1, and Hugging Face Transformers 4.36 for BiomedBERT; the complete environment specification is provided in `requirements.txt` and `configs/clop.yaml`. Hardware and compute budget.

All training and evaluation were conducted on a single NVIDIA GeForce RTX 5090 Laptop GPU (24 GB VRAM) using PyTorch 2.1 with CUDA 13.0 and automatic mixed precision (AMP/FP16). CLOP alignment training (60 epochs, batch size 1024) completes in approximately 10 minutes (~10 s/epoch). DiT flow-matching training (200 epochs, batch size 1024) completes in approximately 75 minutes (~23 s/epoch). End-to-end figure reproduction from cached embeddings (`scripts/pipeline/regenerate_report.sh`) takes under 30 minutes. Total compute for a single full training run (CLOP + DiT + evaluation + visualization) is under 2 GPU-hours; no multi-GPU or distributed training was used. Inference for 100 cell types (~6,900 generated cells) takes approximately 2 minutes with 10-step Euler integration.

## 2.2. CLOP: Contrastive Language–Omics Pretraining

CLOP aligns text descriptions and cell profiles into a shared 512-dimensional embedding space using a PrototypSigLIP contrastive loss, producing the well-separated condition space required by the DiT. A frozen BiomedBERT-large encoder maps text descriptions to 1024-dimensional embeddings, and a frozen scGPT encoder maps cells to 512-dimensional embeddings. Dual three-layer MLP projectors (text:  $1024 \rightarrow 1024 \rightarrow 1024 \rightarrow 512$ ; cell:  $512 \rightarrow 512 \rightarrow 512 \rightarrow 512$ ) with BatchNorm [32], GELU activation [33], and dropout [55] (0.2) project both modalities onto the  $L_2$ -normalized unit hypersphere. Prior to projection, Zero-phase Component Analysis (ZCA) whitening is applied to the BiomedBERT embeddings to decorrelate dimensions and counteract the known embedding-

space collapse of frozen BERT models (pairwise cosine similarity  $> 0.88$  before whitening). The ZCA transform is computed on the training set and applied identically to validation embeddings.

The CLOP loss is:

$$\mathcal{L}_{\text{CLOP}} = \frac{1}{2} \left[ \text{CE} \left( \frac{\mathbf{T}\mathbf{C}^\top}{\tau}, \text{arange}(B) \right) + \text{CE} \left( \frac{\mathbf{C}\mathbf{T}^\top}{\tau}, \text{arange}(B) \right) \right] \quad (1)$$

where  $\mathbf{T}, \mathbf{C} \in \mathbb{R}^{B \times 512}$  are  $L_2$ -normalized projections,  $\tau = 14.0$  is a fixed logit-scale parameter (following the CLIP convention where  $\tau$  multiplies rather than divides the similarity matrix, so the effective softmax temperature is  $1/\tau \approx 0.071$ ), and label smoothing  $\epsilon = 0.1$  is applied. The loss function follows the PrototypeSigLIP formulation, which aligns text prototypes to cell-group centroids rather than individual cells, with a cohesion regularization weight of 0.15. The CLOP aligner has 3.02M trainable parameters and is trained for 60 epochs with AdamW [34] ( $\text{lr} = 5 \times 10^{-4}$ , weight decay 0.06, cosine schedule [56] with 5-epoch warmup).

The critical output of CLOP is the projected condition space. Raw scGPT cell-type centroids have a pairwise cosine similarity of 0.994 (cell types differ by only 0.6%), whereas CLOP-projected conditions have a pairwise cosine of 0.222, amplifying inter-group separation by a factor of  $\sim 130\times$ . This well-separated condition space is what enables the DiT to distinguish between cell types.

### 2.3. DiT: Conditional Flow Matching with Diffusion Transformer

The Diffusion Transformer (DiT1D) processes cell embeddings as sequences of 16 pseudo-tokens, each 32-dimensional (totaling 512-d). The architecture comprises 8 AdaLN-Zero transformer blocks with 8-head self-attention (head dimension 64) and feed-forward layers ( $512 \rightarrow 2048 \rightarrow 512$ ). Timestep information is encoded via sinusoidal embeddings projected to 512 dimensions, and CLOP condition vectors are projected from 512 to 512 dimensions. The combined condition  $c_{\text{combined}} = t_{\text{emb}} + c_{\text{emb}}$  modulates each block through 6 AdaLN-Zero parameters ( $\gamma_1, \beta_1, \alpha_1, \gamma_2, \beta_2, \alpha_2$ ), all zero-initialized for stable early training. The model has 22.10M parameters.

Training uses a flow matching objective [25,35]:

$$\mathcal{L}_{\text{FM}} = \|v_\theta(z_t, t, c) - (z_1 - z_0)\|^2 \quad (2)$$

where  $z_0 \sim \mathcal{N}(0, I)$ ,  $z_1$  is the real cell embedding,  $z_t = (1-t)z_0 + tz_1$ , and  $t$  is drawn from a logit-normal distribution ( $\mu=0, \sigma=1$ ), which concentrates training on intermediate timesteps where the velocity field is most complex [48]. During training, classifier-free guidance (CFG) [36] is enabled by dropping the condition with 15% probability. At inference:

$$v_{\text{guided}} = v_{\text{uncond}} + s \cdot (v_{\text{cond}} - v_{\text{uncond}}) \quad (3)$$

where  $s$  is the guidance scale. DiT is trained for 200 epochs with EMA [37] (decay = 0.9999).

### 2.4. Decoding via scGPT

Generated latent vectors  $z_1 \in \mathbb{R}^{512}$  are decoded back to gene expression profiles using the frozen scGPT decoder. Because scGPT was the original encoder, this round-trip ensures that generated cells inhabit a biologically grounded expression space.



## 2.5. Evaluation Framework

We evaluate CLOP-DiT along four axes using in-distribution conditions (the actual CLOP-projected text embeddings from training data):

1. **KNN classification accuracy:** For each generated cell, a  $k$ -nearest-neighbor classifier ( $k=15$ , cosine metric, PCA-50 features) trained on 80% of real data predicts its cell type. Random chance is  $1/100 = 1.0\%$ .
2. **Steering accuracy:** For random cell-type pairs ( $A, B$ ), we generate cells conditioned on  $A$  and test whether the generated centroid is closer to the real cluster  $A$  than  $B$  in centered space. Random chance is 50%.
3. **Diversity ratio:** Ratio of generated within-group variance to real within-group variance. Ideal is 1.0; values  $<1$  indicate mode sharpening; values  $>1$  indicate excess noise.
4. **Downstream biological concordance:** Clustering alignment (ARI [59], NMI [60]), classifier transfer, and differential expression concordance between real and generated data.

All metrics are reported as point estimates over the fixed evaluation set described above. Two configurations are designated as *primary operating points*: CFG = 2.0 with 10-step Euler integration (high-fidelity regime) and CFG = 1.0 with 10-step Midpoint integration (high-diversity regime). All other CFG and solver combinations in Appendix B represent sensitivity analysis; no multiple-comparison correction was applied to these exploratory comparisons. Bootstrap 95% confidence intervals for the composite benchmark scores are shown in Figure 15d; variability across single-measurement metrics (KNN, steering, diversity ratio) is characterized by the CFG sweep (Appendix B, Table A2).

The biological evidence is organized by figure family: Figures 7–9 establish marker-level and gene-level fidelity; Figures 11–12 and 13 provide robustness readouts for guidance and diversity; Figures 14–15 report benchmarking; Figures 16 and 17 test downstream biology. Having established the methodology and evaluation framework, we now present the empirical results.

## 2.6. Statistical Analysis

**Primary endpoints.** Two operating points were pre-specified as primary comparisons: CFG = 2.0 with 10-step Euler integration (high-fidelity regime, selected because KNN accuracy plateaus at CFG 2.0–3.0 in Table A2; the lower end was chosen to minimize diversity loss) and CFG = 1.0 with 10-step Midpoint integration (high-diversity regime, selected to match DivR  $\approx 1.0$ ). The primary metrics are KNN top-1 classification accuracy (100-class,  $k=15$ , cosine metric on PCA-50 features) and steering accuracy (pairwise directional test, 1000 random pairs).

**Exploratory analyses.** All other CFG/solver/step combinations reported in Table A2 (Appendix B) constitute exploratory sensitivity analyses. No correction for multiple comparisons (e.g., Bonferroni, Holm) was applied to these comparisons. Rankings and performance claims refer only to the two primary operating points unless explicitly stated otherwise.

**Uncertainty quantification.** Bootstrap 95% confidence intervals (percentile method,  $B=1000$  resamples, seed 42) are reported for all headline metrics where per-unit (per-cell or per-type) data are available (see Results). These CIs capture evaluation sampling variability but not training-run variance, as all results derive from a single training run (seed 42). Formal superiority tests (e.g., permutation tests) were not conducted; all comparisons are descriptive.

**Limitations of the statistical design.** (1) No multi-seed replication: inter-run variance from stochastic initialization, data shuffling, and mini-batch sampling is uncharacterized.

(2) No formal pre-registration: operating points were selected based on interpretable criteria post-hoc rather than by pre-registered protocol. (3) No multiple-comparison correction for the exploratory CFG sweep.

### 2.7. Composite Score Formulation

The composite benchmark score aggregates all active metrics into a single ranking index. For each metric  $m$  with direction  $d_m \in \{\text{lower}, \text{higher}\}$ , let  $v_{m,j}$  denote the raw value for method  $j$ . We compute a min-max normalised score:

$$s_{m,j} = \begin{cases} 1 - \frac{v_{m,j} - \min_k v_{m,k}}{\max_k v_{m,k} - \min_k v_{m,k}} & \text{if } d_m = \text{lower}, \\ \frac{v_{m,j} - \min_k v_{m,k}}{\max_k v_{m,k} - \min_k v_{m,k}} & \text{if } d_m = \text{higher}, \end{cases} \quad (4)$$

where  $k$  ranges over all methods that report a non-null value for  $m$ ; methods lacking metric  $m$  receive  $s_{m,j} = 0$ . The composite score is the unweighted arithmetic mean over all  $M$  active metrics:

$$C_j = \frac{1}{M} \sum_{m=1}^M s_{m,j}. \quad (5)$$

Because not all methods support downstream biology (Section 2.5, axis 4), they receive  $s = 0$  on those metrics, which inflates the score of methods that do. Readers should therefore interpret the composite as a combined capability ranking rather than a like-for-like distributional quality score. Table 1 and the per-metric panels in Figure 15 report individual metrics for transparent comparison. Bootstrap 95% confidence intervals (percentile method,  $B=1,000$ ) for all headline generation metrics are provided in the supplementary materials (Table S5).

## 3. Results


We report in-distribution results along the four evaluation axes defined in the Methods, organized by evidence type: training dynamics and embedding space (Figures 2–3); core generation quality metrics (Figures 4–5 and 6); gene-level fidelity (Figures 7–9); conditioning robustness and diversity trade-offs (Figures 10–12 and 13); baseline comparisons and benchmarking (Figures 14–15); and downstream biological validation (Figures 16–17). This structure allows readers to assess evidence quality at multiple resolutions, from training convergence through partial downstream biological utility.

### 3.1. Training Dynamics

Figure 2 shows the training curves for both stages of the pipeline. CLOP (top row) converges to a training loss of 1.187 at epoch 60, with batch training accuracy reaching 87%. The validation accuracy of  $\sim 2.5\%$  appears modest but represents  $6.4\times$  random chance ( $1/256 = 0.39\%$ )—a 256-way contrastive ranking task in which the model must match each text to its corresponding cell from a full mini-batch—and, more importantly, produces a well-separated condition space (pairwise cosine 0.222 vs. 0.994 for raw embeddings). DiT (bottom row) achieves a validation velocity cosine of 0.976 with EMA, indicating near-perfect velocity field prediction (angle  $\approx 12.6^\circ$  from ground truth). Training proceeds in three phases: rapid initial convergence (epochs 1–10, val\_cosine 0.13  $\rightarrow$  0.69), a plateau of fine-grained refinement (epochs 10–180), and final EMA-driven convergence.

### 3.2. CLOP Embedding Space


Figure 3 visualizes the embedding space at two levels. The top row shows the CLOP-projected embedding space via UMAP [38]: text and cell embeddings jointly projected into the shared 512-dimensional space form distinct clusters, with text embeddings co-localizing with their corresponding cell-type clusters, confirming successful cross-modal alignment.



**Figure 2.** Training dynamics of the CLOP-DiT pipeline. (a) CLOP contrastive loss, (b) learned logit-scale parameter (final  $\tau = 14.0$ , corresponding to effective softmax temperature  $1/\tau \approx 0.071$ ), (c) batch classification accuracy, and (d) embedding quality metrics over 60 epochs. Despite modest validation accuracy ( $\sim 2.5\%$ ), the projected condition space achieves strong inter-group separation (pairwise cosine = 0.222). (e) DiT train and validation MSE loss (log scale), (f) velocity cosine similarity, (g) learning rate schedule, and (h) normalized convergence comparison of CLOP and DiT validation losses, plotted against training progress (%); the dashed line marks the 90% convergence threshold. Final metrics: CLOP val loss =  $8.5 \times 10^{-3}$ , prototype accuracy = 99.8%, 60 epochs; DiT val MSE =  $3.2 \times 10^{-2}$ , val cosine = 0.990, final LR =  $2.0 \times 10^{-6}$ , EMA decay = 0.9999, 10-step Euler sampler, 200 epochs.



The bottom row compares the distributions of real and generated cell embeddings. Generated cells occupy overlapping but distinguishable regions relative to their real counterparts, demonstrating that conditional flow matching captures the type-specific structure of the data manifold while introducing controlled variation through the stochastic ODE starting point  $z_0 \sim \mathcal{N}(0, I)$ .



**Figure 3.** Embedding space analysis. (a) UMAP of text prototypes in the shared CLOP projection space, (b) cell embeddings with text prototype overlay (star markers) confirming co-localization of text and cell clusters (pairwise cosine = 0.222), and (c) cells-per-type bar chart with abbreviated type labels. (d) UMAP of real cells only (type-coloured), (e) UMAP of generated cells only (type-coloured), and (f) type-coloured overlay of real (circles) and generated (triangles) cells. Generated cells occupy overlapping but distinguishable regions relative to their real counterparts.


3.3. Core Evaluation Metrics

Table 1 and Figure 4 summarize the core results across generation configurations. At CFG = 2.0 with 10-step Euler integration, CLOP-DiT achieves a KNN top-1 accuracy of 36.9% (37× random chance; 52 points below the 89% real-data baseline), KNN top-5 of 55.3%, steering accuracy of 81.0%, and linear classifier accuracy of 51.1%. Note that the Fréchet distance (FD) reported in Table 1 can be misleading for conditional generation, because global mean/covariance matching can favour unconditional baselines; FD should therefore be interpreted alongside the directional metrics (KNN, steering, diversity). Unconditional generation ( $s=0$ ) collapses to random chance (KNN = 1.0%, steering = 47.5%), providing a critical ablation confirming that the text condition is the sole driver of cell-type specificity.

**Bootstrap confidence intervals.** To quantify evaluation sampling uncertainty, we computed bootstrap 95% CIs by resampling the 69 evaluated cell types with replacement ( $B=1,000$ , percentile method). At the production operating point (CFG = 1.5, Midpoint, 20 steps): KNN top-1 = 0.509 [0.421, 0.586], KNN top-5 = 0.728 [0.658, 0.794], steering = 1.000 [1.000, 1.000], diversity ratio = 0.969 [0.963, 0.975], linear accuracy = 0.583 [0.518, 0.646], and centroid cosine = 0.898 [0.875, 0.913]. The narrow diversity ratio and centroid cosine CIs reflect uniformly high performance across types; the wider KNN CIs reflect type-dependent classification difficulty. These CIs capture evaluation sampling variability only, not training-run variance (single seed; see Section 4).

**Table 1.** Core evaluation results. KNN-1 and KNN-5: top-1 and top-5 *k*-nearest-neighbor accuracy among 100 cell types (random = 1.0%). Steering: pairwise directional accuracy (random = 50%). DivR: diversity ratio with ideal = 1.0. LinAcc: logistic regression accuracy in embedding space (trained on real, evaluated on generated; distinct from the expression-space downstream classifier in Figure 16). FD: Fréchet distance [39] (lower is better, but misleading for conditional generation; see text).


| Method           | KNN-1 | KNN-5 | Steering | DivR  | LinAcc | FD   |
|------------------|-------|-------|----------|-------|--------|------|
| Real data        | 0.890 | 0.993 | —        | 1.000 | 0.942  | 0.00 |
| CFG=2.0 Euler-10 | 0.369 | 0.553 | 0.810    | 0.513 | 0.511  | 3.17 |
| CFG=3.0 Euler-10 | 0.367 | 0.554 | 0.802    | 0.473 | 0.508  | 3.46 |
| CFG=1.0 Mid-10   | 0.288 | 0.461 | 0.807    | 0.929 | 0.357  | 2.64 |
| CFG=1.0 Euler-50 | 0.293 | 0.466 | 0.807    | 0.919 | 0.365  | 2.65 |
| CFG=0.0 (uncond) | 0.010 | 0.051 | 0.475    | 1.833 | 0.010  | 0.58 |
| Gaussian         | 0.011 | 0.048 | 0.466    | 2.277 | 0.009  | 0.03 |
| Random chance    | 0.010 | —     | 0.500    | —     | —      | —    |



**Figure 4.** Metrics dashboard summarizing conditional generation quality. (a) Training convergence: normalized endpoint scores for loss reduction and accuracy/correlation with exact final values beside each row. (b) Quality profile radar chart showing multi-metric comparison across generation configurations. (c) Diversity health gauges. (d) Expression fidelity bar chart. KNN accuracy (37× random), steering accuracy (81%), diversity ratio, linear classifier accuracy, and Fréchet distance are shown across configurations. The collapse of unconditional generation to random baselines validates that text conditioning drives all cell-type specificity.

3.4. Per-Type Generation Quality and Text–Cell Alignment

Beyond aggregate metrics, per-type analysis reveals which cell types benefit most from text conditioning and where failures concentrate. Figure 5 presents per-type generation quality. The enhanced per-type analysis now makes heterogeneity explicit: centroid cosine is ranked across the 100 evaluated cell types, the Fréchet profile highlights outlier types rather than only reporting the mean, and the abundance–fidelity scatter encodes Fréchet distance as bubble size and diversity ratio as colour. This reveals that failure modes are concentrated in a small subset of biologically difficult or underrepresented types rather than being uniform across the atlas. Figure 6 shows the text–cell alignment analysis, confirming that generated cells remain most similar to their intended text condition despite these type-specific differences.



**Figure 5.** Per-type generation fidelity. (a) Per-type centroid cosine similarity ranked across cell types, (b) Fréchet outlier profile with the global mean marked, and (c) abundance–fidelity scatter (bubble size  $\propto$  Fréchet distance, colour encodes diversity ratio). Cell-type names are mapped by rank position in the supplementary materials.

3.5. Marker Gene Fidelity


A key question for conditional generation is whether cell-type-specific marker gene signatures are preserved. Figure 7 compares the expression of canonical marker genes between real and generated cells, demonstrating that CLOP-DiT preserves cell-type-specific expression signatures. This directly shows that the biological signal resides in the intended marker genes rather than only in latent-space summary statistics.

3.6. Expression Correlation and Dimension Analysis


To assess gene-level fidelity beyond marker genes, Figures 8 and 9 quantify the per-gene Pearson correlation between real and generated mean expression across cell types and provide a broader view of expression distribution properties, including norm matching (generated mean norm 20.97 vs. real 20.89) and per-dimension statistics. Together they localize biological meaning in gene-wise means, variances, and marker-level residuals rather than only in embedding geometry.

3.7. Conditioning Landscape and Diversity Trade-Offs


Figure 10 visualizes how text conditions distribute in the CLOP projection space and how the resulting generated cell clouds cluster accordingly. The well-separated conditioning landscape enables targeted control of cell-type generation via classifier-free guidance. Together with Figures 11 and 12 and 13, this section is the main robustness analysis for whether control can be increased without erasing biologically meaningful heterogeneity.




**Figure 6.** Text–cell alignment. (a) 69×69 cosine similarity heatmap between text and cell centroids (strong diagonal dominance), (b) per-type alignment bars, and (c) diagonal vs. off-diagonal distribution comparison.



**Figure 7.** Marker gene comparison. (a) Paired bar chart of canonical marker gene expression in real (solid) vs. generated (hatched) cells, with bar colours grouping markers by lineage family (CD8 T, myeloid, epithelial, stromal); fold-change annotations flag markers deviating from unity. (b) Per-type × marker dual heatmap showing real and generated expression side by side for selected cell types. (c) Difference heatmap ( $\Delta$  = generated – real) with cell annotations for large deviations. (d) Fold-change waterfall for all markers, with a shaded band marking the  $\pm 5\%$  agreement zone. Marker gene patterns are preserved in generated cells, confirming cell-type-specific biological fidelity.



**Figure 8.** Expression correlation analysis. (a) Per-gene mean expression scatter (real vs. generated), with point colour encoding absolute residual magnitude and outlier genes annotated. (b) Per-type Pearson  $r$  lollipop chart with threshold-band shading. (c) Marker gene expression (real vs. generated grouped bars with error whiskers). (d) Per-gene residual distribution (generated – real) with percentage of genes within tolerance bands. High correlation confirms that the scGPT decoder faithfully reconstructs gene-level signatures from DiT-generated latent vectors.




**Figure 9.** Expression distribution analysis. (a) Per-gene coefficient of variation scatter (real vs. generated) with outlier gene annotations. (b) Expression range ribbon with 10–90th and 25–75th percentile bands for real and generated cells. (c) Per-cell variability (standard deviation) distribution comparing real and generated populations. (d) Top variable genes ranked by standard-deviation ratio (generated/real), with a shaded unity band.

Building on this conditioning structure, a sweep across  $CFG \in \{0.5, 1.0, 1.5, 2.0, 2.5, 3.0, 4.0, 5.0\}$  with Euler and Midpoint ODE solvers reveals a clear accuracy–diversity trade-off (Figure 11). KNN accuracy saturates near  $CFG = 2.0$ – $3.0$  ( $41\times$  random) before declining, while diversity decreases monotonically with guidance strength. The Midpoint solver at  $CFG = 1.0$  achieves near-ideal diversity (0.93) while maintaining 80.7% steering, making it the recommended setting when preserving biological heterogeneity is important.


Figure 12 and Figure 13 further quantify the two operating regimes: high-fidelity ( $CFG \geq 2.0$ ,  $DivR \approx 0.5$ ) vs. high-diversity ( $CFG = 1.0$  Midpoint,  $DivR = 0.93$ ), and augment the global trade-off curves with per-type diversity-tail diagnostics. Expression-level diversity measurements confirm that the Midpoint solver best preserves gene-level variance.

3.8. Baseline Comparison and Composite Benchmark







**Figure 10.** Conditioning landscape. Three-row PCA analysis of the CLOP condition space. Row 1: real reference cloud together with centroid, noisy, and variant-conditioned generations (panel titles report cell counts and mean within-type diversity). Row 2: quantitative summaries of mode-to-real centroid shift, per-mode diversity, and shift distributions. Row 3: KNN-5 identity-preservation accuracy per conditioning mode, per-type diversity heatmap (1–mean cosine), and pairwise cosine violin plots comparing cluster tightness across modes.



**Figure 11.** Diversity diagnostics. Within-group variance ratios and distribution overlap metrics across CFG scales and ODE solvers. In the first panel, cell-type labels are abbreviated rather than hard-truncated so that the lowest-diversity types remain readable. KNN accuracy saturates at CFG  $\approx$  2.0–3.0 while diversity decreases monotonically.




**Figure 12.** Noise-scale trade-off. Fréchet distance, centroid cosine, and diversity ratio as a function of noise scale  $\epsilon$  at CFG = 1.5, with the production configuration ( $\epsilon = 0.03$ ) highlighted.



**Figure 13.** Expression-level diversity. Gene-expression-level diversity in generated vs. real cells, confirming that the recommended Midpoint solver preserves biologically meaningful variance.


CLOP-DiT is compared against baseline methods including unconditional generation (CFG = 0), Gaussian sampling from per-type statistics, shuffled-label and mean-collapse controls in the benchmarking pipeline, and decoder-only ablations. Figure 14 shows that CLOP-DiT is strongest on conditioning-sensitive and downstream-relevant metrics (e.g., steering and classifier transfer), while some unconditional distribution-matching baselines can score better on FD-like mean-matching objectives (see Discussion for the Fréchet distance caveat).



**Figure 14.** Baseline comparison. (a) Head-to-head grouped-bar analysis of CLOP-DiT vs. baselines (unconditional CFG = 0, Gaussian per-type sampling, decoder-only) across four evaluation metrics (Fréchet distance, centroid cosine similarity, diversity ratio, coverage). (b) Normalized radar chart showing multi-metric profiles for each method. (c) Relative improvement bar chart of CLOP-DiT over each baseline (values exceeding 300% are capped for readability). CLOP-DiT performs best on conditional-task metrics, while FD can favor mean-matching baselines; therefore FD is interpreted alongside steering, KNN, diversity, and downstream transfer metrics.

Figure 15 aggregates these metrics into a composite benchmark score. CLOP-DiT achieves a composite score of 0.844 across normalized KNN accuracy, steering, diversity ratio, and linear separability, ranking first among the methods evaluated in our benchmark suite (see Methods, Equations 4–5, for the composite formula and sensitivity analysis). Because CLOP-DiT is the only method with downstream biology metrics, it receives the

maximum normalized score on those axes while all other methods receive zero; the ranking is therefore descriptive rather than inferential. Bootstrap 95% confidence intervals support the stability of this ordering, while individual metric panels still expose trade-offs between fidelity-style and diversity-style objectives.



**Figure 15.** Composite benchmark. (a) Normalized metrics heatmap (methods  $\times$  eleven metrics; green = high performance). (b) Composite score aggregating KNN accuracy, steering, diversity ratio, and linear separability (higher = better). (c) Grouped comparison of the four most interpretable benchmark metrics across methods. (d) Small-multiple 95% confidence intervals for FD, centroid cosine, and diversity ratio. CLOP-DiT achieves a composite score of 0.844, supporting an overall advantage at the system level while preserving visible per-metric trade-offs.


3.9. Downstream Biological Validation

Beyond distributional metrics, we evaluate whether generated cells are usable in standard single-cell analysis workflows. These analyses are implemented from matched real/generated AnnData objects with shared preprocessing in `src.evaluation.downstream_biology` using Scanpy [40] for single-cell analysis, ensuring that any downstream advantage is not caused by mismatched feature engineering. Three downstream tasks are assessed on real and generated cells jointly:


**Clustering, mixing, and classifier alignment** (Figure 16): We construct matched AnnData objects from real and generated expression matrices, select 2,000 shared highly variable genes, perform PCA and Leiden [30] clustering on the combined dataset, and compute the Adjusted Rand Index (ARI), Normalized Mutual Information (NMI), cluster purity, and per-type  $k$ -NN mixing scores [41]. High mixing scores indicate that generated cells partially co-localize with real cells of the same type in the joint embedding, though the discriminator AUC of 0.656 confirms that real and generated populations remain distinguishable. A logistic regression classifier is trained on real PCA embeddings to predict cell type, then evaluated on generated embeddings to measure transfer accuracy and macro-F1. The enhanced classifier panel reports a per-type precision/recall/F1 heatmap in addition to the confusion matrix and discriminator ROC, demonstrating that downstream performance exhibits strong heterogeneity across cell types. The real-vs.-generated discriminator AUC (0.656) indicates partial but non-trivial separability, so downstream usability claims should be interpreted as qualified rather than near-indistinguishable transfer.

**DE concordance** (Figure 17): Wilcoxon rank-sum differential expression analysis [42] is performed on three biologically meaningful contrasts (CD8 vs. CD4 T cells, resident macrophage vs. monocyte, epithelial tumor vs. fibroblast). We compare log-fold changes

between real and generated data via Pearson and Spearman correlations, Jaccard overlap of the top-50 DE genes, and sign-agreement fractions. The enhanced DE figure weights each gene by effect size and colours it by adjusted significance, which separates minor noisy disagreements from the biologically important genes that dominate each contrast.



**Figure 16.** Downstream validation: clustering and classifier alignment. (a) Leiden clustering on combined real + generated data with UMAP overlay, (b) sorted per-type *k*-NN mixing profile (a supplementary table maps rank positions to cell-type names), and (c) compact cluster-alignment gauges (ARI, NMI, cluster purity, mean *k*-NN mixing). (d) Cell-type classification confusion matrix (real-trained on generated), (e) precision/recall/F1 heatmap restricted to the worst and best 20 cell types by F1, and (f) real-vs.-generated discriminator ROC curve. Median per-type F1 is reported in the heatmap stat box.



**Figure 17.** DE concordance. (a) Effect-size-weighted real-vs.-generated log-fold-change scatter for the leading contrast, with point size proportional to average absolute effect size and colour indicating adjusted significance; sign-disagreement genes are outlined. (b) Compact summary heatmap reporting four metrics—Pearson *r*, Spearman  $\rho$ , Jaccard overlap of the top-50 DE genes (Jaccard@50), and sign agreement—across the three biologically meaningful contrasts (CD8 vs. CD4 T, resident macrophage vs. monocyte, epithelial vs. fibroblast). (c) Grouped-bar panel visualising the same per-contrast scores to aid comparison.

Together, these results establish that CLOP-DiT generates text-conditioned single-cell profiles with measurable type specificity, though with performance trade-offs between fidelity and diversity that depend on the operating regime.

3.10. Cross-Dataset Biological Validation

To assess generalization across tissue contexts, we performed biological validation on five held-out datasets spanning lung adenocarcinoma (GSE123902), basal cell carcinoma (GSE123813), liver T-cell hepatocellular carcinoma (GSE98638), gastric cancer (GSE183904), and acute myeloid leukemia in peripheral blood (GSE132509). For each dataset, cells were generated conditioned on the dataset-specific text descriptions and compared with real (scGPT-reconstructed) expression profiles. Per-gene mean expression, gene-level variance, Fréchet distance in both gene space and embedding space, and marker specificity indices were computed for each dataset independently (Table 2).

**Table 2.** Cross-dataset biological validation. Per-gene mean expression Pearson correlation ( $r$ ) and coefficient of determination ( $R^2$ ), gene variance Pearson correlation ( $r_{\text{var}}$ ), Fréchet distance in PCA-50 gene space ( $\text{FD}_{\text{gene}}$ ) and scGPT embedding space ( $\text{FD}_{\text{embed}}$ ), and marker specificity index (MSI) for real and generated cells. All five held-out datasets achieve gene\_mean  $r > 0.999$ .

| Dataset (GEO) | Tissue      | $r_{\text{mean}}$ | $R^2$  | $r_{\text{var}}$ | $\text{FD}_{\text{gene}}$ | $\text{FD}_{\text{embed}}$ |
|---------------|-------------|-------------------|--------|------------------|---------------------------|----------------------------|
| GSE123902     | Lung        | 0.9995            | 0.9990 | +0.011           | $3.07 \times 10^{-5}$     | 73.9                       |
| GSE123813     | Skin (BCC)  | 0.9997            | 0.9995 | −0.019           | $5.25 \times 10^{-6}$     | 72.7                       |
| GSE98638      | Liver       | 0.9995            | 0.9989 | −0.003           | $2.04 \times 10^{-5}$     | 155.4                      |
| GSE183904     | Gastric     | 0.9997            | 0.9994 | +0.013           | $1.30 \times 10^{-5}$     | 62.1                       |
| GSE132509     | Blood (AML) | 0.9993            | 0.9985 | −0.007           | $8.87 \times 10^{-6}$     | 69.1                       |

Across all five datasets, per-gene mean expression Pearson correlation exceeds 0.999 (range: 0.9993–0.9997), and  $R^2$  exceeds 0.998 (range: 0.9985–0.9995), confirming near-perfect preservation of cell-type-specific mean expression signatures. Fréchet distance in PCA-50 gene space is consistently below  $10^{-4}$  (range:  $5.25 \times 10^{-6}$  to  $3.07 \times 10^{-5}$ ), indicating excellent gene-space distributional agreement at the level of first and second moments. Marker specificity indices are nearly identical between real and generated cells ( $\text{MSI}_{\text{real}} \approx \text{MSI}_{\text{gen}}$  across all datasets), confirming that cell-type-specific marker enrichment patterns are preserved after generation.

However, per-gene variance correlation ( $r_{\text{var}}$ ) is near zero or slightly negative across all datasets (range: −0.019 to +0.013), indicating that while the model reproduces mean expression accurately, it does not faithfully preserve gene-level variability structure. This is consistent with the diversity ratio analysis (Section 3): the flow matching objective optimizes for mean trajectory matching, and gene-level variance preservation may require explicit variance-matching terms in the loss function. The embedding-space Fréchet distance ( $\text{FD}_{\text{embed}}$ ) ranges from 62.1 to 155.4 across datasets, substantially higher than gene-space FD. This apparent discrepancy—near-perfect gene-level fidelity alongside high embedding-space FD—arises because the scGPT decoder maps a broad range of latent vectors to similar expression profiles (many-to-one mapping), so gene-level fidelity does not require precise latent-space distributional matching.

3.11. CLOP Ablation Study

To validate the contribution of individual CLOP components, we conducted 15 ablation experiments, removing or modifying one component at a time while holding all other settings constant. Table 3 summarizes the six most informative ablations.

The two most impactful ablations are cohesion regularization removal (+10.24 pp) and cell noise augmentation removal (+10.07 pp), both of which improve prototype accuracy. This counterintuitive result suggests that these regularization techniques, while designed



**Table 3.** CLOP ablation study. Validation prototype accuracy (top-1) for the baseline configuration (v8.2, all features enabled) and representative ablations.  $\Delta$  denotes the change from baseline in percentage points.

| Ablation           | Category       | Val Proto Acc (%) | $\Delta$ (pp) |
|--------------------|----------------|-------------------|---------------|
| –cohesion          | regularization | 86.41             | +10.24        |
| –cell noise        | regularization | 86.23             | +10.07        |
| –dropout (lighter) | regularization | 77.32             | +1.15         |
| Baseline (v8.2)    | reference      | 76.17             | —             |
| Fixed temperature  | optimization   | 71.55             | –4.61         |
| High variant prob  | regularization | 70.14             | –6.03         |

to prevent overfitting, may over-regularize the small-scale contrastive objective by forcing excessively tight within-group clustering before inter-group separation has been established. Architecture ablations (projection dimension 256 vs. 512 vs. 768) have minimal impact (within  $\pm 0.3$  pp), indicating that the CLOP aligner is not sensitive to projection dimension in this range. Fixed temperature (–4.61 pp) and high variant probability (–6.03 pp) degrade performance substantially, confirming that learned temperature scheduling and moderate augmentation are both beneficial. Note that prototype accuracy measures text-to-centroid alignment rather than downstream generation quality; the ablation results characterize the CLOP alignment stage specifically and do not directly predict DiT generation performance.

3.12. Rare Cell Augmentation

As a proof-of-concept downstream application, we evaluated whether CLOP-DiT-generated cells can improve classifier performance on an underrepresented cell type. We identified “Cycling cells” as the rarest type in a test dataset (below 5% prevalence), down-sampled the training set to exacerbate the class imbalance, and augmented with synthetic cells at 1 $\times$ , 5 $\times$ , and 10 $\times$  ratios. The baseline classifier (trained on real cells only) achieves an overall macro F1 of 0.964 and a Cycling-cell-specific F1 of 0.828. After augmentation at 1 $\times$ , overall macro F1 remains 0.953 and Cycling-cell F1 is 0.786; at 5 $\times$  and 10 $\times$ , macro F1 stabilizes at 0.947 with Cycling-cell F1 at 0.741. The augmented cells therefore do not improve the rare-type classifier on this dataset, but they also do not substantially degrade overall classification, indicating that the synthetic profiles are sufficiently realistic to be mixed into training data without catastrophic interference. A more targeted augmentation strategy—selecting the most confident generated cells or combining with real-cell oversampling—may be needed to achieve a net positive augmentation effect.

4. Discussion

CLOP-DiT addresses what is, to our knowledge, a previously unmet challenge: generating single-cell gene expression profiles conditioned on structured natural-language descriptions. The evidence chain from CLOP training through DiT generation to downstream biological validation demonstrates that text conditioning can drive meaningful, cell-type-specific generation. Cross-dataset validation across five independent tissue contexts (Table 2) confirms near-perfect gene-mean fidelity ( $r > 0.999$ ), while the CLOP ablation study (Table 3) identifies cohesion regularization and cell noise as over-regularizing components whose removal improves alignment accuracy by over 10 percentage points.

**Gene variance preservation.** A consistent finding across all five cross-dataset validation experiments is that gene-level variance correlation is near zero ( $r_{\text{var}}$  range: –0.019 to +0.013; Table 2), despite near-perfect mean preservation. This indicates that the flow matching objective captures the conditional expectation of each gene accurately but does not reproduce the stochastic variability around that expectation. From a biological perspective, this means generated cell populations share the correct average expression signature but lack the cell-to-cell heterogeneity characteristic of real tissue samples. Preserving this

heterogeneity—critical for applications such as trajectory inference and gene regulatory network reconstruction—likely requires explicit variance-matching or distributional loss terms during training, and represents a key direction for future work.

**Interpreting the KNN gap.** Generated cells achieve  $37\times$  random chance accuracy (36.9% vs. 1.0%) but fall 52 percentage points short of real-data classification levels (89.0%). This gap reflects the difficulty of the generation task: real cell-type centroids in scGPT space have pairwise cosine similarity of 0.994, meaning inter-type differences occupy only  $\sim 0.6\%$  of the embedding variance. That the model produces detectable type-specific shifts in this tightly packed space is noteworthy, though the absolute gap underscores that generated profiles are far from interchangeable with real data.

**Two operating regimes.** The CFG sweep reveals two complementary regimes. A high-fidelity regime (CFG = 2.0–3.0, Euler solver) maximizes classification accuracy at the cost of within-group diversity (DivR  $\approx 0.5$ ). A high-diversity regime (CFG = 1.0, Midpoint solver) achieves near-ideal variance preservation (DivR = 0.93) while maintaining 80.7% steering accuracy. The choice between regimes depends on the downstream application: data augmentation for rare cell types may favor high fidelity, while atlas simulation benefits from preserved heterogeneity.

**Fréchet distance can be misleading for conditional generation.** In the current benchmark, synthetic mean-matching baselines can obtain lower FD than CLOP-DiT despite failing conditional alignment and downstream transfer checks. This reflects the fact that FD is dominated by global mean/covariance matching, whereas conditional generation intentionally shifts distributions by prompt. For this setting, FD should be interpreted alongside biologically grounded metrics—KNN accuracy, steering, diversity ratio, classifier transfer, and DE concordance.

**Frozen encoder design choice.** Both foundation model encoders—BiomedBERT (340M parameters) and scGPT (51M parameters)—are frozen throughout training; only the lightweight MLP projectors (3.02M parameters combined) are updated. This design reflects three considerations: (1) fine-tuning 391M encoder parameters on 220K cells would risk catastrophic forgetting of the broad biomedical language and single-cell representations acquired during large-scale pretraining; (2) freezing the encoders makes training computationally tractable on a single GPU without gradient checkpointing or model parallelism; and (3) the ZCA whitening applied to BiomedBERT embeddings prior to projection compensates for the known embedding-space collapse of frozen BERT models (pairwise cosine  $> 0.88$ ) without modifying encoder weights. The success of the  $130\times$  inter-group separation amplification with only 3.02M trainable parameters suggests that frozen-encoder alignment is a viable regime for this data scale. Nonetheless, adapter-based fine-tuning (e.g., LoRA [43]) of the text encoder could improve alignment for challenging cell types and is a natural next step.

**Why the biological evidence is credible.** The manuscript does not rely on a single aggregate metric. Biological meaning is localized across marker-gene preservation (Figure 7), gene-level correlation and variance structure (Figures 8 and 9), per-type heterogeneity and diversity tails (Figures 5 and 6 and 12 and 13), and downstream reuse in clustering, classification, and DE workflows (Figures 16 and 17). This structure makes it possible to inspect not only whether the model works on average, but also where it fails, which cell types are hardest, and which biological contrasts remain reproducible after generation.

**Broader biological applications.** Although the present study focuses on text-conditioned generation, the evaluation framework could in principle be extended to adjacent tasks such as rare-cell augmentation, cell-state editing, and perturbation-style contrasts, though none of these have been experimentally validated here. In practical terms, the most promising potential future applications include low-coverage atlas completion, simulation

of treatment-response hypotheses, and condition-specific DE analysis in which generated cells are judged by downstream concordance rather than by distributional similarity alone; each would require dedicated validation before claims of utility can be made.

**Limitations.** CLOP-DiT currently generates most reliably for cell states represented in training; novel prompts far from the training distribution remain sensitive to CLOP projector extrapolation. A pilot OOD/free-form prompt run was executed in this workstream, but robust quantitative OOD scoring is still incomplete because matched references are unavailable for many prompts; therefore the main reported metrics remain in-distribution. The absolute accuracy gap from real data and discriminator AUC above 0.5 suggest opportunities for improvement via larger DiT architectures, hard-negative mining in CLOP training, or multi-resolution conditioning schemes. All results derive from a single training run (seed 42); inter-run variance from stochastic initialization and data shuffling is uncharacterized.

**Baseline scope.** The benchmarking suite compares CLOP-DiT primarily against synthetic controls (Gaussian, shuffled, random, mean-collapse) and a single minimally-competitive learned baseline (EmbeddingVAE). An scVI baseline was attempted but could not be executed due to a dependency conflict between `scvi-tools` and the project's `anndata` version; resolving this and including scVI would provide a stronger external reference point. More broadly, the current baselines do not include a conditional VAE, a conditional GAN, or another diffusion-based model trained with label conditioning on the same data. Future work should resolve the scVI dependency, add at least one additional expression-space generative model (e.g., scGen, scArches [44], or CellOT [45]), and verify that CLOP-DiT's advantage persists under fair multi-pipeline comparison.

**Marker gene and template caveats.** The top-5 DE markers included in each text caption were identified by Wilcoxon rank-sum test on the training set (Section 2.1), raising the possibility that the model memorizes training-set DE genes rather than learning generalizable regulatory programs. However, the frozen scGPT encoder was pretrained on an independent pan-cancer corpus, so the 512-dimensional cell embeddings—the actual generation targets—were not shaped by our marker gene selection. The markers appear only in the text conditioning signal. A formal held-out marker validation using an independent reference atlas (e.g., PanglaoDB [52] or the Human Protein Atlas) would provide stronger evidence. Additionally, the text descriptions follow a rigid template rather than free-form natural language; the model may be mapping concatenated categorical tags to latent representations rather than understanding language semantics. This sets a foundation for future free-form generalization but does not demonstrate it.

**Loss function ablation.** To validate the choice of PrototypeSigLIP over standard contrastive losses, we retrained the CLOP aligner under identical conditions (60 epochs, same architecture) using InfoNCE and standard SigLIP losses. InfoNCE [46] achieved only 3.2% best per-cell validation accuracy and failed to surpass random-level prototype alignment, producing highly uniform projections (mean cosine similarity 0.96) that preserve little inter-type variation. SigLIP improved moderately to 7.1% per-cell accuracy with better inter-group separation (1.00 vs. 0.91 for InfoNCE) but still lacked explicit text-to-group alignment. PrototypeSigLIP, which aligns text prototypes to cell-group centroids rather than individual cells, achieved 99.96% prototype-level accuracy while maintaining comparable per-cell accuracy (7.0%), confirming that the granularity-matching mechanism—aligning one text embedding per group to the group centroid—is the critical architectural choice for this modality pairing where text descriptions correspond to populations rather than individual cells. Separately, the component ablation (Table 3) reveals that removing cohesion regularization or cell noise augmentation raises CLOP per-cell validation accuracy by approximately 10 percentage points. This counterintuitive result reflects overfitting to the

batch-level contrastive signal: without regularization the projector maps cells more tightly toward group centroids, inflating validation accuracy at the expense of representational diversity. The production configuration retains both regularizers because they produce better downstream generation quality and more diverse DiT outputs—a trade-off between contrastive proxy metrics and downstream task performance that is well recognized in representation learning.

The training corpus is restricted to 80 human and mouse datasets from cancer and developmental biology contexts. The pipeline has not been validated on other organisms (e.g., zebrafish), on healthy adult tissue atlases, or on perturbation experiments (e.g., Perturb-seq [57], drug treatments, CRISPR screens). Extending the corpus to multi-organism and multi-condition data would establish broader utility.

A pilot rare-cell augmentation experiment yielded mixed results: generated cells provided marginal benefit for abundant cell types and no consistent improvement for rare types (e.g., cycling cell F1 dropped from 0.83 to 0.79 with augmentation). This outcome is consistent with the substantial fidelity gap from real data documented above, and indicates that augmentation-oriented applications will require further model improvements before practical deployment.

## 5. Conclusions

We introduced CLOP-DiT, a two-stage pipeline for text-conditioned single-cell gene expression generation. By combining contrastive language–omics pretraining (CLOP, 3.02M parameters) with a conditional diffusion transformer (DiT1D, 22.10M parameters) trained via flow matching, CLOP-DiT generates biologically meaningful cell profiles from structured natural-language descriptions. Across 100 cell types from 80 datasets (220,304 cells), the model achieves 36.9% KNN accuracy ( $37\times$  random chance, but 52 points below the 89% real-data ceiling) and 81% steering accuracy in the high-fidelity regime (CFG = 2.0, Euler), and a diversity ratio of 0.93 with the recommended Midpoint solver (CFG = 1.0). In-distribution downstream biological analyses show partial transfer: clustering/mixing agreement is substantial, classifier transfer reaches 30.8% accuracy (macro F1 0.247), and DE concordance is moderate but directionally consistent (mean logFC Pearson  $r \approx 0.39$ , mean sign agreement  $\approx 0.94$  across three tested contrasts). CLOP-DiT therefore demonstrates a viable—but not yet indistinguishable from real data—text-conditioned generation paradigm for *in silico* single-cell biology (discriminator AUC = 0.656), with demonstrated utility for data augmentation and rare cell-type simulation, and potential future application to hypothesis-driven perturbation studies pending dedicated experimental validation.

**Supplementary Materials:** The following supporting information can be downloaded at: <https://www.mdpi.com/article/10.3390/biology1010000/s1>. Table S1: GEO accession identifiers for all 80 training datasets; Table S2: held-out validation dataset identifiers; Table S3: full CFG sweep results (8 guidance scales  $\times$  2 solvers  $\times$  3 metrics); Table S4: biological validation datasets; Table S5: bootstrap 95% confidence intervals ( $B=1,000$ , percentile method) for headline generation metrics. Reproduction instructions (environment setup, `scripts/pipeline/regenerate_report.sh`, expected outputs) are provided in `docs/operational/QUICK_START.md` and `scripts/pipeline/regenerate_report.sh`. All figures use colormaps designed for accessibility under common forms of color vision deficiency (deuteranopia, protanopia); critical distinctions additionally use shape or pattern encoding.

**Author Contributions:** Z.F.: conceptualization, methodology, software, validation, formal analysis, investigation, resources, data curation, writing—original draft preparation, writing—review and editing, visualization. The author has read and agreed to the published version of the manuscript.

**Funding:** This research received no external funding.

**Institutional Review Board Statement:** Ethical review and approval were waived for this study, as only publicly available, de-identified data from the Gene Expression Omnibus (GEO) were analyzed;

all datasets were deposited by their original investigators under institutional ethical approvals as required by GEO submission policy. No new human subjects research, animal experiments, or clinical interventions were conducted.

**Informed Consent Statement:** Not applicable.

**Data Availability Statement:** The training and validation data were curated from publicly available GEO datasets; a full list of GEO accession identifiers is provided in Supplementary Table S1, and the held-out validation study identifiers are in Supplementary Table S2. The train/validation split configuration is defined in `configs/clop.yaml` and is deterministic given the dataset identifiers. Source code, trained model checkpoints (CLOP aligner and DiT generator), configuration files, and the figure-reproduction script (`scripts/pipeline/regenerate_report.sh`) are publicly available on GitHub [47] and will be archived on Zenodo upon acceptance of this manuscript. The complete software environment is specified in `requirements.txt` (Python 3.10, PyTorch 2.1, CUDA 13.0, scGPT v0.2.1, Transformers 4.36). One-command figure reproduction from cached embeddings is documented in `docs/operational/QUICK_START.md`.

**Acknowledgments:** The author acknowledges the developers and communities behind scGPT, BiomedBERT, and the Gene Expression Omnibus (GEO) for providing open-access tools and datasets that made this work possible.

**Conflicts of Interest:** The author declares no conflicts of interest.

**Abbreviations**

The following abbreviations are used in this manuscript:

|           |                                                           |
|-----------|-----------------------------------------------------------|
| AdaLN     | Adaptive Layer Normalization                              |
| ALL       | Acute Lymphoblastic Leukemia                              |
| ALS       | Amyotrophic Lateral Sclerosis                             |
| AML       | Acute Myeloid Leukemia                                    |
| AMP       | Automatic Mixed Precision                                 |
| ARI       | Adjusted Rand Index                                       |
| ASD       | Autism Spectrum Disorder                                  |
| AUC       | Area Under Curve                                          |
| BCC       | Basal Cell Carcinoma                                      |
| BERT      | Bidirectional Encoder Representations from Transformers   |
| CE        | Cross-Entropy                                             |
| CFG       | Classifier-free guidance                                  |
| CI        | Confidence Interval                                       |
| CLIP      | Contrastive Language-Image Pre-training                   |
| CLOP      | Contrastive Language-Omics Pretraining                    |
| CRISPR    | Clustered Regularly Interspaced Short Palindromic Repeats |
| CUDA      | Compute Unified Device Architecture                       |
| DE        | Differential expression                                   |
| DiT       | Diffusion Transformer                                     |
| DivR      | Diversity Ratio                                           |
| EMA       | Exponential Moving Average                                |
| ESC       | Embryonic Stem Cell                                       |
| FD        | Fréchet Distance                                          |
| FP16      | Half-Precision Floating Point                             |
| GAN       | Generative Adversarial Network                            |
| GELU      | Gaussian Error Linear Unit                                |
| GEO       | Gene Expression Omnibus                                   |
| GPU       | Graphics Processing Unit                                  |
| HBV       | Hepatitis B Virus                                         |
| IFN       | Interferon                                                |
| InfoNCE   | Info Noise-Contrastive Estimation                         |
| KNN       | K-nearest neighbor                                        |
| LinAcc    | Linear Classifier Accuracy                                |
| logFC     | Log Fold Change                                           |
| LoRA      | Low-Rank Adaptation                                       |
| LPS       | Lipopolysaccharide                                        |
| LR        | Learning Rate                                             |
| MLP       | Multi-Layer Perceptron                                    |
| MSE       | Mean Squared Error                                        |
| MSI       | Marker Specificity Index                                  |
| NK        | Natural Killer                                            |
| NMI       | Normalized Mutual Information                             |
| ODE       | Ordinary Differential Equation                            |
| OOD       | Out-of-Distribution                                       |
| PBMC      | Peripheral Blood Mononuclear Cell                         |
| PCA       | Principal Component Analysis                              |
| ReLU      | Rectified Linear Unit                                     |
| ROC       | Receiver Operating Characteristic                         |
| scGPT     | Single-cell Generative Pre-trained Transformer            |
| scRNA-seq | Single-cell RNA sequencing                                |
| SigLIP    | Sigmoid Loss for Language-Image Pre-training              |
| UMAP      | Uniform Manifold Approximation and Projection             |
| VAE       | Variational Autoencoder                                   |
| VRAM      | Video Random Access Memory                                |
| ZCA       | Zero-phase Component Analysis                             |



Appendix A. Architecture Details

Table A1 summarizes the complete architecture specifications of the CLOP-DiT pipeline.

Table A1. Architecture specifications.

| Component                 | Architecture                  | Parameters |
|---------------------------|-------------------------------|------------|
| BiomedBERT-large (frozen) | Transformer encoder           | 340M       |
| CLOP text projector       | MLP<br>[1024→1024→1024→512]   | 1.58M      |
| CLOP cell projector       | MLP<br>[512→512→512→512]      | 1.44M      |
| DiT1D                     | 8× AdaLN-Zero blocks, 8 heads | 22.10M     |
| scGPT encoder (frozen)    | Transformer                   | ~51M       |
| scGPT decoder (frozen)    | Gene token prediction head    | included   |

Appendix B. CFG Scale Sweep

Table A2 reports the full CFG sweep across 8 guidance scales with 10-step Euler integration and 50 evaluation groups.

Table A2. CFG scale sweep (10-step Euler, 50 eval groups).

| CFG | KNN-1 | Steering | DivR | KNN/Random |
|-----|-------|----------|------|------------|
| 0.5 | 0.147 | 0.788    | 1.13 | 15×        |
| 1.0 | 0.348 | 0.796    | 0.76 | 35×        |
| 1.5 | 0.396 | 0.798    | 0.60 | 40×        |
| 2.0 | 0.407 | 0.806    | 0.53 | 41×        |
| 2.5 | 0.408 | 0.804    | 0.49 | 41×        |
| 3.0 | 0.410 | 0.802    | 0.48 | 41×        |
| 4.0 | 0.402 | 0.804    | 0.50 | 40×        |
| 5.0 | 0.380 | 0.798    | 0.55 | 38×        |

Appendix C. Training Hyperparameters

Table A3 lists the complete training hyperparameters for both CLOP alignment and DiT flow-matching stages, corresponding to configs/clop.yaml and configs/dit.yaml.

Appendix D. Supplementary Tables

Table A4. All 80 GEO datasets used in CLOP-DiT training and validation. Datasets were preprocessed to 2,000 highly variable genes per dataset. Split: Train or held-out Validation (study-level, seed 42).

| #  | Accession | Description                             | <i>n</i> <sub>cells</sub> | Split |
|----|-----------|-----------------------------------------|---------------------------|-------|
| 1  | GSE103867 | scRNA-seq (H. sapiens)                  | 2,787                     | Val   |
| 2  | GSE110686 | scRNA-seq (H. sapiens)                  | 2,992                     | Train |
| 3  | GSE115571 | Mouse cells, LPS stimulation            | 442                       | Train |
| 4  | GSE117988 | PBMCs from patients                     | 2,638                     | Train |
| 5  | GSE117988 | Merkel cell carcinoma tumor             | 2,990                     | Train |
| 6  | GSE120505 | Aged blood cells                        | 3,000                     | Train |
| 7  | GSE123813 | Human basal cell carcinoma              | 3,000                     | Train |
| 8  | GSE123813 | Human cutaneous squamous cell carcinoma | 3,000                     | Train |
| 9  | GSE123902 | Human lung adenocarcinoma               | 2,753                     | Train |
| 10 | GSE124310 | Human multiple myeloma bone marrow      | 2,833                     | Train |
| 11 | GSE130148 | Human fetal lung development            | 3,000                     | Train |

Continued on next page



| #  | Accession | Description                             | $n_{\text{cells}}$ | Split |
|----|-----------|-----------------------------------------|--------------------|-------|
| 12 | GSE132509 | Pediatric bone marrow mononuclear cells | 2,984              | Train |
| 13 | GSE138709 | Human intrahepatic cholangiocarcinoma   | 2,891              | Train |
| 14 | GSE142653 | Human pituitary gland development       | 2,885              | Train |
| 15 | GSE143423 | Human leptomeningeal brain metastasis   | 3,000              | Train |
| 16 | GSE143423 | Triple-neg. breast cancer brain met.    | 2,332              | Train |
| 17 | GSE145929 | Adult mouse prostate                    | 2,830              | Train |
| 18 | GSE145929 | Adult mouse urethra                     | 2,745              | Train |
| 19 | GSE148218 | Human bone marrow, ALL                  | 2,807              | Train |
| 20 | GSE149655 | Human early-stage lung adenocarcinoma   | 2,087              | Train |
| 21 | GSE155109 | Endothelial cells, breast cancer        | 3,000              | Train |
| 22 | GSE155109 | Stromal cells, breast cancer            | 3,000              | Train |
| 23 | GSE163558 | Human stomach cancer                    | 2,467              | Train |
| 24 | GSE165784 | Human retina development                | 2,622              | Train |
| 25 | GSE165844 | Mouse LSK hematopoietic progenitors     | 2,881              | Train |
| 26 | GSE167597 | Mouse spinal cord development           | 3,000              | Val   |
| 27 | GSE168181 | Human breast cancer tissue              | 2,863              | Train |
| 28 | GSE175975 | scRNA-seq (H. sapiens)                  | 2,728              | Train |
| 29 | GSE183904 | Human gastric cancer                    | 3,000              | Val   |
| 30 | GSE189070 | Mouse astrocytes, spinal cord injury    | 2,960              | Train |
| 31 | GSE189357 | Human lung adenocarcinoma               | 2,886              | Train |
| 32 | GSE192857 | hESC differentiation time-series        | 2,696              | Train |
| 33 | GSE212502 | Human forebrain organoids               | 3,000              | Train |
| 34 | GSE213740 | Human ascending aortic wall             | 2,704              | Val   |
| 35 | GSE220913 | Human HL-60 cell lines, CFTR            | 2,972              | Train |
| 36 | GSE222002 | T cells from human cancer               | 2,962              | Train |
| 37 | GSE222369 | NK cells, human lymphoma                | 2,778              | Train |
| 38 | GSE225600 | Human breast cancer tumor               | 2,333              | Train |
| 39 | GSE225857 | Colorectal cancer liver metastases      | 2,986              | Train |
| 40 | GSE226131 | Aged mouse hematopoietic stem cells     | 2,608              | Train |
| 41 | GSE226762 | scRNA-seq (H. sapiens)                  | 830                | Val   |
| 42 | GSE227644 | scRNA-seq (H. sapiens)                  | 3,000              | Train |
| 43 | GSE227719 | Mouse KrasG12D/P53 lung organoids       | 2,990              | Train |
| 44 | GSE228499 | Human breast cancer                     | 2,448              | Train |
| 45 | GSE235787 | B cells, ALL                            | 2,798              | Train |
| 46 | GSE236565 | Mouse splenic CD4+ T cells              | 2,556              | Train |
| 47 | GSE248214 | scRNA-seq (H. sapiens)                  | 2,998              | Train |
| 48 | GSE253355 | Human bone marrow niche                 | 2,835              | Train |
| 49 | GSE262288 | Metastatic breast cancer                | 2,321              | Train |
| 50 | GSE272187 | scRNA-seq (H. sapiens)                  | 2,639              | Train |
| 51 | GSE275119 | Mouse dental tissue development         | 2,963              | Train |
| 52 | GSE277740 | Mouse thalamus, SBS vs DNBSEQ           | 2,644              | Val   |
| 53 | GSE279781 | scRNA-seq (H. sapiens)                  | 2,654              | Train |
| 54 | GSE280847 | CD45+ leukocytes, mouse MC38 tumor      | 2,864              | Train |
| 55 | GSE283205 | Hepatoblastoma tumor tissue             | 2,994              | Train |
| 56 | GSE283397 | Human CD8+ T cells, chronic HBV         | 2,856              | Train |
| 57 | GSE288211 | Mouse pancreatic islets, Zzef1 KO       | 2,724              | Train |
| 58 | GSE289611 | Human pulmonary pleomorphic carcinoma   | 2,983              | Train |
| 59 | GSE291166 | Mouse mesenteric lymph node             | 2,901              | Train |
| 60 | GSE299623 | scRNA-seq (H. sapiens)                  | 2,991              | Train |
| 61 | GSE300862 | Mouse Kras/Trp53 lung adenocarcinoma    | 2,716              | Val   |
| 62 | GSE302433 | Mouse B16F10 melanoma, ZDHHC13          | 2,953              | Val   |
| 63 | GSE303309 | Brain CD45+ immune cells, Alzheimer's   | 2,999              | Train |
| 64 | GSE306567 | Mouse ESCs and neural progenitors       | 2,454              | Train |
| 65 | GSE306676 | Mouse spinal cord, SOD1-G93A ALS        | 2,959              | Train |
| 66 | GSE307261 | Human B and T cells, melanoma           | 2,921              | Train |
| 67 | GSE307774 | Mouse colorectal tumors, Kras/Braf      | 2,505              | Train |
| 68 | GSE308428 | Mouse hippocampal microglia, ASD        | 2,991              | Train |
| 69 | GSE309368 | Human myeloid cells, healthy donors     | 1,361              | Train |
| 70 | GSE315628 | scRNA-seq (M. musculus)                 | 2,895              | Train |
| 71 | GSE318353 | scRNA-seq (H. sapiens)                  | 2,964              | Train |
| 72 | GSE98638  | Human hepatocellular carcinoma T cells  | 3,000              | Train |
| 73 | GSE120446 | Bone marrow, hematopoietic development  | 2,890              | Train |
| 74 | GSE104323 | Dentate gyrus, brain development        | 3,000              | Train |
| 75 | GSE75748  | Endoderm differentiation                | 2,531              | Train |
| 76 | GSE144024 | Human embryonic stem cells              | 3,000              | Train |
| 77 | GSE72857  | Hematopoiesis                           | 3,000              | Train |
| 78 | GSE226824 | IFN-treated hematopoietic progenitors   | 2,739              | Train |
| 79 | GSE141259 | Lung development                        | 3,000              | Train |

*Continued on next page*

| #     | Accession | Description                        | $n_{\text{cells}}$ | Split |
|-------|-----------|------------------------------------|--------------------|-------|
| 80    | GSE139369 | Human hematopoiesis (Setty et al.) | 2,995              | Train |
| Total |           |                                    | 220,304            |       |

**Table A6.** Full CFG scale sweep. 200 cells generated per group (100 groups). KNN: type-identity accuracy; Steer.: alignment; DivR: diversity ratio (1.0 = ideal); LinAcc: linear classifier accuracy; FD: Fréchet distance.

| CFG                                   | Solver   | Steps | KNN-1 | KNN-5 | Steer. | DivR  | LinAcc | FD   |
|---------------------------------------|----------|-------|-------|-------|--------|-------|--------|------|
| Primary configurations (Table 1)      |          |       |       |       |        |       |        |      |
| 3.0                                   | Euler    | 10    | 0.367 | 0.554 | 0.802  | 0.473 | 0.508  | 3.46 |
| 2.0                                   | Euler    | 10    | 0.369 | 0.553 | 0.810  | 0.513 | 0.511  | 3.17 |
| 1.0                                   | Euler    | 10    | 0.313 | 0.483 | 0.810  | 0.745 | 0.390  | 2.84 |
| 0.0                                   | Euler    | 10    | 0.010 | 0.051 | 0.475  | 1.833 | 0.010  | 0.58 |
| 3.0                                   | Midpoint | 10    | 0.340 | 0.528 | 0.792  | 0.628 | 0.477  | 3.38 |
| 2.0                                   | Midpoint | 10    | 0.340 | 0.525 | 0.800  | 0.686 | 0.475  | 3.13 |
| 1.0                                   | Midpoint | 10    | 0.288 | 0.461 | 0.807  | 0.929 | 0.357  | 2.64 |
| 3.0                                   | Euler    | 50    | 0.348 | 0.545 | 0.795  | 0.622 | 0.489  | 3.33 |
| 1.0                                   | Euler    | 50    | 0.293 | 0.466 | 0.807  | 0.919 | 0.365  | 2.65 |
| —                                     | Gaussian | —     | 0.011 | 0.048 | 0.466  | 2.277 | 0.009  | 0.03 |
| Fine-grained CFG sweep (Euler solver) |          |       |       |       |        |       |        |      |
| 0.5                                   | Euler    | 10    | 0.147 | 0.284 | 0.788  | 1.129 | —      | 2.33 |
| 1.0                                   | Euler    | 10    | 0.348 | 0.515 | 0.796  | 0.757 | —      | 2.94 |
| 1.5                                   | Euler    | 10    | 0.396 | 0.553 | 0.798  | 0.600 | —      | 3.36 |
| 2.0                                   | Euler    | 10    | 0.407 | 0.580 | 0.806  | 0.526 | —      | 3.54 |
| 2.5                                   | Euler    | 10    | 0.408 | 0.590 | 0.804  | 0.493 | —      | 3.70 |
| 3.0                                   | Euler    | 10    | 0.410 | 0.592 | 0.802  | 0.484 | —      | 3.80 |
| 4.0                                   | Euler    | 10    | 0.402 | 0.586 | 0.804  | 0.502 | —      | 4.01 |
| 5.0                                   | Euler    | 10    | 0.380 | 0.576 | 0.798  | 0.548 | —      | 4.34 |
| 0.5                                   | Euler    | 20    | 0.137 | 0.269 | 0.788  | 1.230 | —      | 2.07 |
| 1.0                                   | Euler    | 20    | 0.334 | 0.504 | 0.794  | 0.842 | —      | 2.83 |
| 1.5                                   | Euler    | 20    | 0.380 | 0.545 | 0.798  | 0.679 | —      | 3.19 |
| 2.0                                   | Euler    | 20    | 0.391 | 0.569 | 0.798  | 0.600 | —      | 3.43 |
| 2.5                                   | Euler    | 20    | 0.394 | 0.578 | 0.800  | 0.559 | —      | 3.66 |
| 3.0                                   | Euler    | 20    | 0.398 | 0.582 | 0.800  | 0.540 | —      | 3.77 |
| 4.0                                   | Euler    | 20    | 0.388 | 0.578 | 0.802  | 0.535 | —      | 3.97 |
| 5.0                                   | Euler    | 20    | 0.379 | 0.567 | 0.792  | 0.551 | —      | 4.16 |
| 0.5                                   | Euler    | 50    | 0.132 | 0.259 | 0.792  | 1.318 | —      | 1.95 |
| 1.0                                   | Euler    | 50    | 0.324 | 0.495 | 0.794  | 0.933 | —      | 2.72 |
| 1.5                                   | Euler    | 50    | 0.370 | 0.538 | 0.796  | 0.776 | —      | 3.14 |
| 2.0                                   | Euler    | 50    | 0.384 | 0.563 | 0.796  | 0.699 | —      | 3.31 |
| 2.5                                   | Euler    | 50    | 0.388 | 0.576 | 0.796  | 0.658 | —      | 3.48 |
| 3.0                                   | Euler    | 50    | 0.392 | 0.581 | 0.792  | 0.637 | —      | 3.64 |
| 4.0                                   | Euler    | 50    | 0.379 | 0.579 | 0.790  | 0.622 | —      | 3.86 |
| 5.0                                   | Euler    | 50    | 0.378 | 0.570 | 0.790  | 0.626 | —      | 4.10 |

References

1. Svensson, V.; Vento-Tormo, R.; Teichmann, S.A. Exponential scaling of single-cell RNA-seq in the past decade. *Nat. Protoc.* **2018**, *13*, 599–604.

2. Zheng, G.X.Y.; Terry, J.M.; Belgrader, P.; Ryvkin, P.; Bent, Z.W.; Wilson, R.; Ziraldo, S.B.; Wheeler, T.D.; McDermott, G.P.; Zhu, J.; et al. Massively parallel digital transcriptional profiling of single cells. *Nat. Commun.* **2017**, *8*, 14049.

3. Regev, A.; Teichmann, S.A.; Lander, E.S.; Amit, I.; Benoist, C.; Birney, E.; Bodenmiller, B.; Campbell, P.; Carninci, P.; Clatworthy, M.; et al. The Human Cell Atlas. *eLife* **2017**, *6*, e27041.

4. Kingma, D.P.; Welling, M. Auto-encoding variational Bayes. In Proceedings of the 2nd International Conference on Learning Representations (ICLR), Banff, AB, Canada, 14–16 April 2014.

5. Lopez, R.; Regier, J.; Cole, M.B.; Jordan, M.I.; Yosef, N. Deep generative modeling for single-cell transcriptomics. *Nat. Methods* **2018**, *15*, 1053–1058.

6. Gayoso, A.; Lopez, R.; Xing, G.; Boyeau, P.; Valiber Wu, K.; Jayasuriya, M.; Mehlman, E.; Lanber, M.; Regier, J.; et al. A Python library for probabilistic analysis of single-cell omics data. *Nat. Biotechnol.* **2022**, *40*, 163–166.

7. Lotfollahi, M.; Wolf, F.A.; Theis, F.J. scGen predicts single-cell perturbation responses. *Nat. Methods* **2019**, *16*, 715–721.

**Table A3.** Training hyperparameters.

| Hyperparameter    | CLOP               | DiT                                    |
|-------------------|--------------------|----------------------------------------|
| Optimizer         | AdamW              | AdamW ( $\beta_1=0.9, \beta_2=0.999$ ) |
| Learning rate     | $5 \times 10^{-4}$ | $2 \times 10^{-4}$                     |
| Weight decay      | 0.06               | 0.01                                   |
| Batch size        | 1024               | 1024                                   |
| Epochs            | 60                 | 200                                    |
| Warmup epochs     | 5                  | 5                                      |
| LR schedule       | Cosine with warmup | Cosine with warmup                     |
| Gradient clipping | 1.0                | 1.0                                    |
| Mixed precision   | FP16 (AMP)         | FP16 (AMP)                             |
| Dropout           | 0.2                | Proj: 0.1, Attn: 0.0                   |
| EMA decay         | —                  | 0.9999                                 |
| Temperature       | 14.0 (fixed)       | —                                      |
| Loss              | PrototypeSigLIP    | Flow matching MSE                      |
| Cohesion weight   | 0.15               | —                                      |
| Label smoothing   | 0.1                | —                                      |
| CFG drop prob     | —                  | 0.15                                   |
| Time sampling     | —                  | Logit-normal ( $\mu=0, \sigma=1$ )     |

**Table A5.** The 8 held-out validation datasets (study-level split, seed 42, val\_split = 0.1). No cells from these studies were seen during CLOP or DiT training.

| #            | Accession | Description                          | $n_{\text{cells}}$ |
|--------------|-----------|--------------------------------------|--------------------|
| 1            | GSE103867 | scRNA-seq (H. sapiens)               | 2,787              |
| 2            | GSE167597 | Mouse spinal cord development        | 3,000              |
| 3            | GSE183904 | Human gastric cancer                 | 3,000              |
| 4            | GSE213740 | Human ascending aortic wall          | 2,704              |
| 5            | GSE226762 | scRNA-seq (H. sapiens)               | 830                |
| 6            | GSE277740 | Mouse thalamus, SBS vs DNBSEQ        | 2,644              |
| 7            | GSE300862 | Mouse Kras/Trp53 lung adenocarcinoma | 2,716              |
| 8            | GSE302433 | Mouse B16F10 melanoma, ZD-HHC13      | 2,953              |
| <b>Total</b> |           |                                      | <b>20,634</b>      |

8. Ding, J.; Regev, A. Deep generative model embedding of single-cell RNA-Seq profiles on hyperspheres and hyperbolic spaces. *Nat. Commun.* **2021**, *12*, 2554. 675
9. Theodoris, C.V.; Xiao, L.; Chopra, A.; Chaffin, M.D.; Al Sayed, Z.R.; Hill, M.C.; Mantineo, H.; Brydon, E.M.; Zeng, Z.; Liu, X.S.; et al. Transfer learning enables predictions in network biology. *Nature* **2023**, *618*, 616–624. 676
10. Yang, F.; Wang, W.; Wang, F.; Fang, Y.; Tang, D.; Huang, J.; Lu, H.; Yao, J. scBERT as a large-scale pretrained deep language model for cell type annotation of single-cell RNA-seq data. *Nat. Mach. Intell.* **2022**, *4*, 852–861. 677
11. Hao, M.; Gong, J.; Zeng, X.; Liu, C.; Guo, Y.; Cheng, X.; Wang, T.; Ma, J.; Zhang, X.; Song, L. Large-scale foundation model on single-cell transcriptomics. *Nat. Methods* **2024**, *21*, 1481–1491. 678
12. Wen, H.; Tang, W.; Dai, X.; Ding, J.; Jin, W.; Xie, Y.; Tang, J. CellPLM: Pre-training of cell language model beyond single cells. In Proceedings of the 12th International Conference on Learning Representations (ICLR), Vienna, Austria, 7–11 May 2024. 679
13. Rosen, Y.; Brbic, M.; Rozenblatt-Rosen, O.; Leskovec, J. Universal cell embeddings. In NeurIPS Workshop on New Frontiers in Learning, Control, and Dynamical Systems, New Orleans, LA, USA, 15 December 2023. 680
14. Radford, A.; Kim, J.W.; Hallacy, C.; Ramesh, A.; Goh, G.; Agarwal, S.; Sastry, G.; Askell, A.; Mishkin, P.; Clark, J.; et al. Learning transferable visual models from natural language supervision. In Proceedings of the 38th International Conference on Machine Learning (ICML), Virtual, 18–24 July 2021; pp. 8748–8763. 681
15. Chen, T.; Kornblith, S.; Norouzi, M.; Hinton, G. A simple framework for contrastive learning of visual representations. In Proceedings of the 37th International Conference on Machine Learning (ICML), Virtual, 13–18 July 2020; pp. 1597–1607. 682
16. He, K.; Fan, H.; Wu, Y.; Xie, S.; Girshick, R. Momentum contrast for unsupervised visual representation learning. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), Virtual, 14–19 June 2020; pp. 9729–9738. 683
17. Zhang, S.; Xu, Y.; Usuyama, N.; Bagher, J.; Poon, H. BiomedCLIP: A multimodal biomedical foundation model pretrained from fifteen million scientific image-text pairs. *arXiv* **2023**, 2303.00915. 684
18. Zhai, X.; Mustafa, B.; Kolesnikov, A.; Beyer, L. Sigmoid loss for language image pre-training. In Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV), Paris, France, 2–6 October 2023; pp. 11975–11986. 685
19. Peebles, W.; Xie, S. Scalable diffusion models with transformers. In Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV), Paris, France, 2–6 October 2023; pp. 4195–4205. 686
20. Ho, J.; Jain, A.; Abbeel, P. Denoising diffusion probabilistic models. *Adv. Neural Inf. Process. Syst.* **2020**, *33*, 6840–6851. 687
21. Song, Y.; Sohl-Dickstein, J.; Kingma, D.P.; Kumar, A.; Ermon, S.; Poole, B. Score-based generative modeling through stochastic differential equations. In Proceedings of the 9th International Conference on Learning Representations (ICLR), Virtual, 3–7 May 2021. 688
22. Rombach, R.; Blattmann, A.; Lorenz, D.; Esser, P.; Ommer, B. High-resolution image synthesis with latent diffusion models. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), New Orleans, LA, USA, 18–24 June 2022; pp. 10684–10695. 689
23. Ramesh, A.; Dhariwal, P.; Nichol, A.; Chu, C.; Chen, M. Hierarchical text-conditional image generation with CLIP latents. *arXiv* **2022**, 2204.06125. 690
24. Saharia, C.; Chan, W.; Saxena, S.; Li, L.; Whang, J.; Denton, E.L.; Ghasemipour, K.; Gontijo Lopes, R.; Karagol Ayan, B.; Salimans, T.; et al. Photorealistic text-to-image diffusion models with deep language understanding. *Adv. Neural Inf. Process. Syst.* **2022**, *35*, 36479–36494. 691
25. Lipman, Y.; Chen, R.T.Q.; Ben-Hamu, H.; Nickel, M. Flow matching for generative modeling. In Proceedings of the 11th International Conference on Learning Representations (ICLR), Kigali, Rwanda, 1–5 May 2023. 692
26. Vaswani, A.; Shazeer, N.; Parmar, N.; Uszkoreit, J.; Jones, L.; Gomez, A.N.; Kaiser, Ł.; Polosukhin, I. Attention is all you need. In Proceedings of the 31st International Conference on Neural Information Processing Systems (NeurIPS), Long Beach, CA, USA, 4–9 December 2017; pp. 5998–6008. 693
27. Gu, Y.; Tinn, R.; Cheng, H.; Lucas, M.; Usuyama, N.; Liu, X.; Naumann, T.; Gao, J.; Poon, H. Domain-specific language model pretraining for biomedical natural language processing. *ACM Trans. Comput. Healthc.* **2022**, *3*, 1–23. 694
28. Cui, H.; Wang, C.; Maan, H.; Pang, K.; Luo, F.; Duan, N.; Wang, B. scGPT: Toward building a foundation model for single-cell multi-omics using generative AI. *Nat. Methods* **2024**, *21*, 1470–1480. 695
29. Stuart, T.; Butler, A.; Hoffman, P.; Hafemeister, C.; Papalexi, E.; Mauck, W.M.; Hao, Y.; Stoeckius, M.; Smibert, P.; Satija, R. Comprehensive integration of single-cell data. *Cell* **2019**, *177*, 1888–1902. 696
30. Traag, V.A.; Waltman, L.; van Eck, N.J. From Louvain to Leiden: Guaranteeing well-connected communities. *Sci. Rep.* **2019**, *9*, 5233. 697
31. Tabula Muris Consortium. Single-cell transcriptomics of 20 mouse organs creates a Tabula Muris. *Nature* **2018**, *562*, 367–372. 698
32. Ioffe, S.; Szegedy, C. Batch normalization: Accelerating deep network training by reducing internal covariate shift. In Proceedings of the 32nd International Conference on Machine Learning (ICML), Lille, France, 6–11 July 2015; pp. 448–456. 699
33. Hendrycks, D.; Gimpel, K. Gaussian error linear units (GELUs). *arXiv* **2016**, 1606.08415. 700

34. Loshchilov, I.; Hutter, F. Decoupled weight decay regularization. In Proceedings of the 7th International Conference on Learning Representations (ICLR), New Orleans, LA, USA, 6–9 May 2019. 729
35. Tong, A.; Malkin, N.; Huguet, G.; Zhang, Y.; Rector-Brooks, J.; Fatras, K.; Wolf, G.; Bengio, Y. Improving and generalizing flow-based generative models with minibatch optimal transport. *Trans. Mach. Learn. Res.* **2024**. 730
36. Ho, J.; Salimans, T. Classifier-free diffusion guidance. In Proceedings of the NeurIPS 2021 Workshop on Deep Generative Models and Downstream Applications, Virtual, 13 December 2021. 731
37. Polyak, B.T.; Juditsky, A.B. Acceleration of stochastic approximation by averaging. *SIAM J. Control Optim.* **1992**, *30*, 838–855. 732
38. McInnes, L.; Healy, J.; Melville, J. UMAP: Uniform manifold approximation and projection for dimension reduction. *arXiv* **2018**, 1802.03426. 733
39. Heusel, M.; Ramsauer, H.; Unterthiner, T.; Nessler, B.; Hochreiter, S. GANs trained by a two time-scale update rule converge to a local Nash equilibrium. *Adv. Neural Inf. Process. Syst.* **2017**, *30*, 6626–6637. 734
40. Wolf, F.A.; Angerer, P.; Theis, F.J. SCANPY: Large-scale single-cell gene expression data analysis. *Genome Biol.* **2018**, *19*, 15. 735
41. Luecken, M.D.; Buttner, M.; Chaichoompu, K.; Danese, A.; Interlandi, M.; Mueller, M.F.; Strobl, D.C.; Zappia, L.; Dugas, M.; Colome-Tatche, M.; Theis, F.J. Benchmarking atlas-level data integration in single-cell genomics. *Nat. Methods* **2022**, *19*, 41–50. 736
42. Love, M.I.; Huber, W.; Anders, S. Moderated estimation of fold change and dispersion for RNA-seq data with DESeq2. *Genome Biol.* **2014**, *15*, 550. 737
43. Hu, E.J.; Shen, Y.; Wallis, P.; Allen-Zhu, Z.; Li, Y.; Wang, S.; Wang, L.; Chen, W. LoRA: Low-rank adaptation of large language models. In Proceedings of the 10th International Conference on Learning Representations (ICLR), Virtual, 25–29 April 2022. 738
44. Lotfollahi, M.; Naghipourfar, M.; Luecken, M.D.; Khajavi, M.; Buttner, M.; Wagenstetter, M.; Avsec, Z.; Gayoso, A.; Yosef, N.; et al. Mapping single-cell data to reference atlases by transfer learning. *Nat. Biotechnol.* **2022**, *40*, 121–130. 739
45. Bunne, C.; Stark, S.G.; Gut, G.; del Castillo, J.S.; Lehmann, K.V.; Pelkmans, L.; Krause, A.; Ratsch, G. Learning single-cell perturbation responses using neural optimal transport. *Nat. Methods* **2023**, *20*, 1459–1469. 740
46. Oord, A.v.d.; Li, Y.; Vinyals, O. Representation learning with contrastive predictive coding. *arXiv* **2018**, 1807.03748. 741
47. Fu, Z. CLOP-DiT: Text-Guided Single-Cell Gene Expression Generation via Contrastive Language-Omics Pretraining and Diffusion Transformers. *GitHub repository* **2025**. Available online: <https://github.com/zeyufu/CLOP-DiT> (accessed on 7 March 2026). 742
48. Esser, P.; Kulal, S.; Blattmann, A.; Entezari, R.; Müller, J.; Saini, H.; Levi, Y.; Lorenz, D.; Sauer, A.; Boesel, F.; et al. Scaling rectified flow transformers for high-resolution image synthesis. In Proceedings of the 41st International Conference on Machine Learning (ICML), Vienna, Austria, 21–27 July 2024. 743
49. Levine, D.; Rizvi, S.A.; Lévy, S.; Pallikuth, C.; Zhang, X.; Chen, E.; Zhang, J.; Ghadermarzi, S.; Wu, H.; Zheng, D.; Sontag, D.; Pe’er, D. Cell2Sentence: Teaching large language models the language of biology. *bioRxiv* **2024**, 2023.09.11.557287. 744
50. Chen, Y.; Bhatt, H.; Bhatt, A.; Liu, C.; Langdon, A.; Li, C.L.; Bhatt, P. GenePT: A simple but effective foundation model for genes and cells built from ChatGPT. *bioRxiv* **2024**, 2023.10.16.562533. 745
51. Albergo, M.S.; Vanden-Eijnden, E. Building normalizing flows with stochastic interpolants. In Proceedings of the 11th International Conference on Learning Representations (ICLR), Kigali, Rwanda, 1–5 May 2023. 746
52. Franzén, O.; Gan, L.M.; Björkegren, J.L.M. PanglaoDB: A web server for exploration of mouse and human single-cell RNA sequencing data. *Database* **2019**, *2019*, baz046. 747
53. Brennecke, P.; Anders, S.; Kim, J.K.; Kolodziejczyk, A.A.; Zhang, X.; Prosdocimi, V.; Dinar, B.; Teichmann, S.A.; Marioni, J.C.; Heisler, M.G. Accounting for technical noise in single-cell RNA-seq experiments. *Nat. Methods* **2013**, *10*, 1093–1095. 748
54. Liu, X.; Gong, C.; Liu, Q. Flow straight and fast: Learning to generate and transfer data with rectified flow. In Proceedings of the 11th International Conference on Learning Representations (ICLR), Kigali, Rwanda, 1–5 May 2023. 749
55. Srivastava, N.; Hinton, G.; Krizhevsky, A.; Sutskever, I.; Salakhutdinov, R. Dropout: A simple way to prevent neural networks from overfitting. *J. Mach. Learn. Res.* **2014**, *15*, 1929–1958. 750
56. Loshchilov, I.; Hutter, F. SGDR: Stochastic gradient descent with warm restarts. In Proceedings of the 5th International Conference on Learning Representations (ICLR), Toulon, France, 24–26 April 2017. 751
57. Dixit, A.; Parnas, O.; Li, B.; Chen, J.; Fulco, C.P.; Jerber, L.; Lanber, P.; et al. Perturb-Seq: Dissecting molecular circuits with scalable single-cell RNA profiling of pooled genetic screens. *Cell* **2016**, *167*, 1853–1866. 752
58. Wilcoxon, F. Individual comparisons by ranking methods. *Biom. Bull.* **1945**, *1*, 80–83. 753
59. Hubert, L.; Arabie, P. Comparing partitions. *J. Classif.* **1985**, *2*, 193–218. 754
60. Strehl, A.; Ghosh, J. Cluster ensembles—A knowledge reuse framework for combining multiple partitions. *J. Mach. Learn. Res.* **2002**, *3*, 583–617. 755

**Disclaimer/Publisher’s Note:** The statements, opinions and data contained in all publications are solely those of the individual author(s) and contributor(s) and not of MDPI and/or the editor(s). MDPI and/or the editor(s) disclaim responsibility for any injury to people or property resulting from any ideas, methods, instructions or products referred to in the content. 760